

TEST DOCUMENT

Microsoft Enterprise AI Services Code of Conduct

Summarize this article for me

This Microsoft Enterprise AI Services Code of Conduct (“Code of Conduct”) defines the requirements that all customers of Microsoft AI Services (as defined in the Product Terms) must adhere to in good faith.

This Code of Conduct unifies and replaces the previous codes for Microsoft Generative AI Services, Azure Face in Foundry Tools, and Azure Speech in Foundry Tools text to speech. This Code of Conduct differs from earlier codes of conduct in structure and phrasing, and it is designed to better align with emerging AI regulations (e.g., EU AI Act).

This Code of Conduct applies in addition to the Microsoft Product Terms, including the Acceptable Use Policy. Any capitalized terms not otherwise defined in this Code of Conduct will have the meaning given to them in the applicable Microsoft agreement(s) under which customers purchase the Online Service.

Responsible AI requirements

Customers must ensure that all of their applications built with Microsoft AI Services, including applications that make decisions, or take actions, autonomously or with varying levels of human intervention:

Implement technical and operational measures to detect fraudulent user behavior in account creation and during use.

Implement strong technical controls on inputs and outputs, including decisions made and actions taken by their applications, to reduce the likelihood of misuse beyond the application's intended purpose.

Disclose when the output, decisions, or actions are generated by AI, including the synthetic nature of generated voices, images, and/or videos, such that users are not likely to be deceived or duped – or able to deceive or dupe others – into believing they are interacting with a real person, that any generated content is authentic or, without their consent, attributable to a specific individual, or that any AI-generated decision or action is done by a human.

Prior to any external display, transmission, or distribution of an AI-generated video, ensure that either (A) your organization has reviewed the video for compliance with this Code of Conduct, or (B) your application that generates the video 1) implements a durable visible watermark reasonably sufficient to inform viewers that the content was generated by AI, 2) updates and re-signs AI Content Credentials (as defined below) as necessary to reflect accurate content provenance, and 3) meets the other requirements of this Code of Conduct. Are tested thoroughly, continuously, and are subject to appropriate human oversight so customers can find and mitigate undesirable behaviors.

Establish feedback channels that allow users to report abuse or issues, and ensure reasonable responses to feedback that is received.

Implement additional scenario-specific mitigations, as appropriate, to ensure responsible use of the Microsoft AI Service, including meaningful human oversight.

Provide all necessary notices and obtain all necessary consents as required by applicable law for both the customer and Microsoft to process data, including third-party data, as part of a customer's use of the Microsoft AI Service

Implement robust security and access control measures, including protecting the Microsoft AI Service resource permissions and having strong user authentication mechanisms.

Usage restrictions

Customers, users, and applications built with Microsoft AI Services must NOT use the services, including applications that make decisions, or take actions, autonomously or with varying levels of human intervention:

In any way that is inconsistent with this Code of Conduct.

In any manner that can inflict harm on individuals, organizations, or society.

To affect individuals in any way that is otherwise prohibited by law or regulation.

To generate, present alongside, monetize, or interact with content, decisions, or actions prohibited in this Code of Conduct.

To make decisions or take actions without appropriate human oversight as part of an application that may have a consequential impact on any individual's legal position, financial position, life opportunities, employment opportunities, or human rights, or may result in physical or psychological harm to an individual.

To deceive or intentionally misinform (e.g., false advertising), or deploy subliminal techniques (e.g., visual, auditory, or other signals beyond a normal person's range of perception) with the intent to manipulate or distort the behavior of a person in a way that causes harm, including through unsolicited phone calls, bulk communications, posts, or messages.

To exploit any of the vulnerabilities of a person due to their age, disability, or a specific socio-economic situation, with the objective, or the effect, of materially distorting the behavior of that person or a person belonging to that group in a manner that causes or is reasonably likely to cause that person or another person significant harm.

For social scoring or predictive profiling that would lead to discriminatory, unfair, biased, detrimental, unfavorable, or harmful treatment of certain persons or groups of persons.

For the assessment of criminality risk of natural persons based solely on the profiling of a natural person or on assessing their personality traits and characteristics. This prohibition shall not apply to AI systems used to support the human assessment of the involvement of a person in a criminal activity, which is already based on objective and verifiable facts directly linked to a criminal activity.

Based on their biometric data, to categorize people or to deduce or infer their race, political opinions, trade union membership, religious or philosophical beliefs, or sex life or sexual orientation, except for any labelling or filtering of lawfully acquired biometric datasets, such as images, based on biometric data or categorizing of biometric data in the area of law enforcement.

To infer people's sensitive attributes (e.g., gender, race, nationality, religion, or specific age), which does not include age range, position of mouth (e.g., smile or frown), and hair color; or infer sensitive information about people without their explicit consent unless used in a lawful manner by a law enforcement, government entity, court, or government official subject to judicial oversight in a jurisdiction that maintains a fair and independent judiciary.

To attempt to infer people's emotional states from their physical, physiological, or behavioral characteristics (e.g., facial expressions, facial movements, or speech patterns), including inferring emotions such as anger, disgust, happiness, sadness, surprise, fear, or other terms commonly used to describe a person's emotional state.

For creating chatbots that (i) are erotic, romantic, or used for erotic or romantic purposes or (ii) are personas of specific people without their explicit consent; or to enable end users to create their own chatbots without oversight.

To create or expand facial recognition databases through the untargeted scraping of facial images from the internet or CCTV footage.

Except to the extent customers are approved by Microsoft for modified content filtering and/or abuse monitoring, to identify or verify individual identities based on people's faces, voices, or other physical, physiological, or behavioral characteristics.

For unlawful tracking, surveillance, stalking, or harassment of a person.

For ongoing surveillance or real-time or near real-time identification or persistent tracking of the individual using any of their personal data, including biometric data, without the individual's valid consent.

For facial recognition purposes (including identification or verification of individual identities) by or for a state or local police department in the United States.

For any real-time facial recognition technology on mobile cameras used by any law enforcement globally to attempt to identify individuals in uncontrolled, "in the wild" environments, which includes (without limitation) police officers on patrol using body-worn or dash-mounted cameras using facial recognition technology to attempt to identify individuals present in a database of suspects or prior inmates.

To detect content credentials or other provenance methods, marks, or signals ("AI Content Credentials") with the purpose of removing or altering them, unless the alterations are to improve the accuracy of the AI Content Credentials (including re-signing AI Content Credentials as needed).

To impersonate any person without explicit and valid consent, including to gain information or privileges, or to simulate the voice or image of politicians or government officials, even with their consent.

Content requirements

Microsoft prohibits the use of Microsoft AI Services for processing, generating, classifying, or filtering content in ways that can inflict harm on individuals, organizations, or society, including but not limited to use of the service for purposes in conflict with this Code of Conduct or the Microsoft Product Terms.

All content released through a customer's use, or integration, of any Microsoft AI Service must be originally created by the publisher, appropriately licensed from the third-party rights holder, used as permitted by the rights holder, or used as otherwise permitted by law. It is the customer's sole responsibility to ensure that customers have appropriate rights to all content input to the Microsoft AI Service (e.g. generated speech and associated metadata).

These content requirements apply to the use of features of, and the output of, all Microsoft AI Services. This includes, but is not limited to, use of features of Azure OpenAI Service and all content provided as input to or generated as output from all models available in Azure OpenAI Service. These requirements also apply to the use of Content Safety in Foundry Control Plane, including features such as custom categories, and to all content provided as input to the service and content generated as output from the service regardless of content filter settings.

Exploitation and abuse

Child sexual exploitation and abuse

Microsoft prohibits content that describes, features, or promotes child sexual exploitation or abuse, whether or not prohibited by law. This includes sexual content involving a child or that sexualizes a child.

Grooming

Microsoft prohibits content that describes or is used for the purposes of grooming children.

Grooming is the act of an adult building a relationship with a child for the purposes of exploitation, especially sexual exploitation. This includes communicating with a child for the purpose of sexual exploitation, trafficking, or other forms of exploitation.

Non-consensual intimate content

Microsoft prohibits content that describes, features, or promotes non-consensual intimate activity whether or not it is sexually explicit.

Sexual solicitation

Microsoft prohibits content that describes, features, promotes, or is used for, purposes of solicitation of commercial sexual activity and sexual services. This includes encouragement and coordination of real sexual activity.